

API Evolution: A Study of Software Stability and Change Trends

Jordan Ell
University of Victoria
British Columbia, Canada
jell@uvic.ca

Braden Simpson
University of Victoria
British Columbia, Canada
braden@uvic.ca

Daniela Damian
University of Victoria
British Columbia, Canada
danielad@csc.uvic.ca

Abstract—Something

I. INTRODUCTION

Intorduction goes here.

II. RELATED WORK

Related work goes here.

III. METHODOLOGY

In order to answer our research question, we decided to use the tool ChangeDistiller created by Fluri et al. [1]. This tool allows us to detect fine grained source code changes in Java projects This tool works by building an abstract syntax tree of a file before and after a code change, then it tries to determine the smallest possible edit distance between the tree. This results in the source code change at a fine grained level performed in the commit.

We took ChangeDistiller and applied it across 10 open source Java projects. For each of the projects, we obtained the software configuration management (SCM) system which is used to store all source code changes of a project. When it was necessary, we converted some forms of SCM system to Git in order to reduce implementation burdens of using multiple SCMs. Once the SCM was obtained, we used ChangeDistiller and iterated over every commit found in a project's git master branch. We stored 34 of ChangeDistiller's built in source code change types for each commit. We noted how many of each change type was performed in each commit and stored that information in a PostgreSQL database. In order to filter and protect our results, we manually inspected the 10 Java projects studied in order to identify code built for test. We separated changes to this test code from all other code to ensure our results only focused on real implementation while allowing us to study changes to test based code separately.

Once the ChangeDistiller information was collected, we decided to examine software change trends surrounding releases of the project's we had selected. Since releases have preconceived notions of software stability, we decided that by studying the types of changes surrounding these releases, we could get a better understanding of what types of source code changes or trends constitute software stability or maturity. In order to study the release points, we went to each of our 10 project's home pages on-line and looked through their release

histories for major, minor, alpha, beta, and release candidate type releases. In total we identified 472 releases across our 10 studied projects.

Once the release dates were collected we set about analyzing out data by creating change ratios surrounding the release dates of each project as a way to measure the trend of a particular change type. To do this we used Equation 1. Equation ?? works by summing a particular change type in commits (T_c) from the release date (r) to a given number of days after the release (d) divided by the number of commits in this date range ($|c|$), then dividing this by the sum of that same particular change type in commits (T_c) from a given number of days (d) before the release date (r) to the release date divided by divided by the number of commits in this date range ($|c|$). This equation gives us a ratio of a particular change type happening before and after a particular release. If the ratio is above 1 then that particular change type occurs more frequently after the release and if it is below 1 then is occurs more frequently before the release. For the purposes of our study, we set the number of days before and after the release (d) to 60 as the projects studies had many months in between their major releases. This quantitative data formed much of the basis for the results to come in Section IV

$$\text{Ratio}_{norm} = \frac{\sum_{c=r}^{r+d} T_c / |c|}{\sum_{c=r}^{r-d} T_c / |c|} \quad (1)$$

Aside from generating quantitative data, we also created a web application for the visualization of the data called API Evolution (APIE). This visualizer allowed us to inspect a single project and a single change type metric at a time (see Figure 1) for qualitative analysis of software evolution trends. We used this tool to manually inspect 4 specific change type trends surround release dates. To do this, we aggregated change types across 50 commits, meaning that each point in the graph represented the date of a commit and the sum of the particular change type's occurrences over the last 50 commits. This was used to smooth out the curves presented by the tool to allow easier manual inspection. Manual inspections were labeled into 4 categories: upward trending, local maximum, downward trending, and local minimum. Since the graphs were quite turbulent, best estimations were given by two judges at each release point to fit the graph into the aforementioned 4




Fig. 1. An screen shot of the APIE visualizer showing project Eclipse.Mylyn.Context with change type PUBLIC_ADDED_METHODS being analyzed.

slope categories. The two judges used 1.5 months before and after the release date as start and end points for the graph trend line.

We performed 1888 manual inspections across 10 projects, 472 release dates and 4 change types, and used this data to form the basis of our qualitative data. Quantitative data was cross validated with qualitative data to form some of the final results to be seen in the next section.

IV. RESULTS

To answer our research question, we conducted a case study of 10 open source Java projects. These projects are: eclipse.jdt.core, eclipse.jdt.ui, eclipse.jetty.project, eclipse.mylyn.common, eclipse.mylyn.context, hibernate-orm, hibernate-ogm, hibernate-search, eclipse.maven.core, and eclipse.maven.surefire. These project were chosen because of their high use amongst other Java projects and to study specific ecosystems of projects and their evolution trends.

Due to the space requirements of this paper, we focus our results on Major releases of the 10 case study projects and a select few of the calculated change ratios. There were 109 major releases across the 10 studied projects. All of the major findings as per values computed from Equation 1 can be seen in Table I.

TABLE I: Implementation oriented change types and their normalized change ratios at 60 days on each side of releases.

Object	Added	Changed	Removed
Public Classes	0.78	0.76	0.67
Public Methods (Signature)	1.12	0.66	2.45
Public Methods (Bodies)	-	0.86	-
Private Classes	0.59	0.92	0.32
Private Methods (Signatures)	0.86	0.91	1.49
Private Methods (Bodies)	-	0.91	-

As it can be seen in Table I, there a few notable change type trends around major releases. We can see that any changes to

public classes are more likely to occur before a major release than after, noting that the addition, changing, and deletion of public classes have change ratios of 0.78, 0.76 and 0.67 respectively. The two most drastic change type trends to note from this table however are the removing of public methods and the removing of private classes. We can see that public methods are more than twice as likely to occur after a major release than before it and the removal of private classes is three times as likely to happen before a major release than after.

In terms of added code, we found as per Table I that additions to private classes are a large indicator of a project being at a stable release point with almost being twice as likely to occur before a major release than after. A large surprise to changed code indicators was that of the public method bodies. We can see that with a change ration of 0.86 there is not a major indication of these types of changes happening before or after a release which leads us to believe that method implementation may not be able to be detected as stable through software metrics. Changes to method signatures also point towards changes before a release rather than after which may indicate that stable software only needs to be stable after the release and not prior, as changing method signature is considered a fairly unstable change [?].

Our qualitative results from manual graph inspections can be seen in Table II. These results show that adding, changing signatures and bodies of, and removing public methods tend to all be at a local minimum of change type trends at major releases.

TABLE II: Qualitative graph analysis results.

Change Type	Upward Trend	Local Maximum
Added Public Methods	21.6%	17.2%
Changed Public Methods (Signature)	6.0%	19.8%
Changed Public Methods (Bodies)	9.2%	16.5%
Removed Public Methods	7.8%	16.4%

Lastly we found that software changes related to testing can be highly used as an indicator of a major release point within the projects studied. The change ratios found can be seen in Table III. As it can be seen, the three largest indicators of a project release and stability with regards to test based changes as the addition of test classes, the removal of test classes, and the removal of test methods.

We summarize our key findings as to source code change type trends which can be found at points of stability as the following: the addition of private classes more often before a stable point, the changing of public method signatures more often before a stable point, the removal of public methods more than twice as often after a stable point, the removal of private classes three times as often after a stable point, and the removal of test classes more than twice as often after a stable point. There are of course other change ratios to consider as listed in this section, however we found the aforementioned to be the most indicative of software stability around major releases.

TABLE III: Test oriented change types and their normalized change ratios at 60 days on each side of releases.

Object	Added	Changed	Removed
Classes	0.72	0.86	0.36
Methods	1.55	0.98	1.93

V. FUTURE WORK

VI. CONCLUSIONS

Software stability is an often used measure in determining some aspects of code quality. This paper has present some initial indications as to how software stability may be found, based on major release stabilities, by analyzing source code change type trends both before and after major software releases. The source code change trends found in this paper, in combination with future work and analysis, will allow for the detection of stable code points throughout a software projects life time. This paper has also shown the beginnings of a visualization for source code change trends which may be used as a visual cue towards project stability and potential areas of instability where action may need to be taken.